

exercise4_109301060

June 4, 2024

1 Exercises for Lecture 4

1.0.1 109301060

1.0.2 Q1

a.

```
[ ]: import statsmodels.formula.api as smf
import scipy.stats as stats
import pandas as pd
import numpy as np

## a.
## Set sample size and number of iterations:
n = 100          ## sample size
r = 50           ## number of iteration

## Set true parameters
beta0 = 0.4
beta1 = -1.46
beta2 = 2.5

df = n - 2 - 1
cv_005 = stats.t.ppf(1 - 0.05, df = df)          ## Critical value for 5% CI
cv_001 = stats.t.ppf(1 - 0.01, df = df)          ## Critical value for 1% CI

## Create empty vectors to store the simulation results
beta0_result_005 = np.empty(r)
beta1_result_005 = np.empty(r)
beta2_result_005 = np.empty(r)

beta0_result_001 = np.empty(r)
beta1_result_001 = np.empty(r)
beta2_result_001 = np.empty(r)

## Set random seed
np.random.seed(1234567)
```

```

## Generate samples of x1 and x2
x1 = stats.uniform.rvs(size = n)
x2 = stats.norm.rvs(loc = 0.78, size = n)

for i in range(r):

    u = np.random.normal(size = n) ## Generate u
    y = beta0 + beta1 * x1 + beta2 * x2 + u ## Generate y
    datax = pd.DataFrame({'y':y, 'x1':x1, 'x2':x2})

    reg = smf.ols(formula = 'y ~ x1 + x2', data = datax)
    results = reg.fit()
    beta_hat = results.params
    bse = results.bse

    ## Whether the true parameters lie in the 95% CI's?
    beta0_result_005[i] = (beta0 >= (beta_hat['Intercept'] -
    ↪cv_005*bse['Intercept'])) and (beta0 <= (beta_hat['Intercept'] +
    ↪cv_005*bse['Intercept']))
    beta1_result_005[i] = (beta1 >= (beta_hat['x1'] - cv_005*bse['x1'])) and
    ↪(beta1 <= (beta_hat['x1'] + cv_005*bse['x1']))
    beta2_result_005[i] = (beta2 >= (beta_hat['x2'] - cv_005*bse['x2'])) and
    ↪(beta2 <= (beta_hat['x2'] + cv_005*bse['x2']))

    ## Whether the true parameters lie in the 99% CI's?
    beta0_result_001[i] = (beta0 >= (beta_hat['Intercept'] -
    ↪cv_001*bse['Intercept'])) and (beta0 <= (beta_hat['Intercept'] +
    ↪cv_001*bse['Intercept']))
    beta1_result_001[i] = (beta1 >= (beta_hat['x1'] - cv_001*bse['x1'])) and
    ↪(beta1 <= (beta_hat['x1'] + cv_001*bse['x1']))
    beta2_result_001[i] = (beta2 >= (beta_hat['x2'] - cv_001*bse['x2'])) and
    ↪(beta2 <= (beta_hat['x2'] + cv_001*bse['x2']))

## Print out the results
## Alpha = 0.05
print(np.mean(beta0_result_005), np.mean(beta1_result_005), np.
    ↪mean(beta2_result_005))
## Alpha = 0.01
print(np.mean(beta0_result_001), np.mean(beta1_result_001), np.
    ↪mean(beta2_result_001))

```

0.94 0.9 0.92
1.0 0.98 0.96

b.

```

[ ]: import statsmodels.formula.api as smf
import scipy.stats as stats
import pandas as pd
import numpy as np

## a.
## Set sample size and number of iterations:
n = 100                ## sample size
r = 200                ## number of iteration

## Set true parameters
beta0 = 0.4
beta1 = -1.46
beta2 = 2.5

df = n - 2 - 1
cv_005 = stats.t.ppf(1 - 0.05, df = df)          ## Critical value for 5% CI
cv_001 = stats.t.ppf(1 - 0.01, df = df)          ## Critical value for 99% CI

## Create empty vectors to store the simulation results
beta0_result_005 = np.empty(r)
beta1_result_005 = np.empty(r)
beta2_result_005 = np.empty(r)

beta0_result_001 = np.empty(r)
beta1_result_001 = np.empty(r)
beta2_result_001 = np.empty(r)

## Set random seed
np.random.seed(1234567)

## Generate samples of x1 and x2
x1 = stats.uniform.rvs(size = n)
x2 = stats.norm.rvs(loc = 0.78, size = n)

for i in range(r):

    u = np.random.normal(size = n)                ## Generate u
    y = beta0 + beta1 * x1 + beta2 * x2 + u        ## Generate y
    datax = pd.DataFrame({'y':y, 'x1':x1, 'x2':x2})

    reg = smf.ols(formula = 'y ~ x1 + x2', data = datax)
    results = reg.fit()
    beta_hat = results.params
    bse = results.bse

```

```

    ## Whether the true parameters lie in the 95% CI's?
    beta0_result_005[i] = (beta0 >= (beta_hat['Intercept'] -
↪cv_005*bse['Intercept'])) and (beta0 <= (beta_hat['Intercept'] +
↪cv_005*bse['Intercept']))
    beta1_result_005[i] = (beta1 >= (beta_hat['x1'] - cv_005*bse['x1'])) and
↪(beta1 <= (beta_hat['x1'] + cv_005*bse['x1']))
    beta2_result_005[i] = (beta2 >= (beta_hat['x2'] - cv_005*bse['x2'])) and
↪(beta2 <= (beta_hat['x2'] + cv_005*bse['x2']))

    ## Whether the true parameters lie in the 99% CI's?
    beta0_result_001[i] = (beta0 >= (beta_hat['Intercept'] -
↪cv_001*bse['Intercept'])) and (beta0 <= (beta_hat['Intercept'] +
↪cv_001*bse['Intercept']))
    beta1_result_001[i] = (beta1 >= (beta_hat['x1'] - cv_001*bse['x1'])) and
↪(beta1 <= (beta_hat['x1'] + cv_001*bse['x1']))
    beta2_result_001[i] = (beta2 >= (beta_hat['x2'] - cv_001*bse['x2'])) and
↪(beta2 <= (beta_hat['x2'] + cv_001*bse['x2']))

## Print out the results
## Alpha = 0.05
print(np.mean(beta0_result_005), np.mean(beta1_result_005), np.
↪mean(beta2_result_005))
## Alpha = 0.01
print(np.mean(beta0_result_001), np.mean(beta1_result_001), np.
↪mean(beta2_result_001))

```

0.895 0.89 0.915

1.0 0.975 0.97

c.

```

[ ]: import statsmodels.formula.api as smf
import scipy.stats as stats
import pandas as pd
import numpy as np

## a.
## Set sample size and number of iterations:
n = 100                ## sample size
r = 1000               ## number of iteration

## Set true parameters
beta0 = 0.4
beta1 = -1.46
beta2 = 2.5

```

```

df = n - 2 - 1
cv_005 = stats.t.ppf(1 - 0.05, df = df)          ## Critical value for 95% CI
cv_001 = stats.t.ppf(1 - 0.01, df = df)          ## Critical value for 99% CI

## Create empty vectors to store the simulation results
beta0_result_005 = np.empty(r)
beta1_result_005 = np.empty(r)
beta2_result_005 = np.empty(r)

beta0_result_001 = np.empty(r)
beta1_result_001 = np.empty(r)
beta2_result_001 = np.empty(r)

## Set random seed
np.random.seed(1234567)

## Generate samples of x1 and x2
x1 = stats.uniform.rvs(size = n)
x2 = stats.norm.rvs(loc = 0.78, size = n)

for i in range(r):

    u = np.random.normal(size = n)                ## Generate u
    y = beta0 + beta1 * x1 + beta2 * x2 + u        ## Generate y
    datax = pd.DataFrame({'y':y, 'x1':x1, 'x2':x2})

    reg = smf.ols(formula = 'y ~ x1 + x2', data = datax)
    results = reg.fit()
    beta_hat = results.params
    bse = results.bse

    ## Whether the true parameters lie in the 95% CI's?
    beta0_result_005[i] = (beta0 >= (beta_hat['Intercept'] -
cv_005*bse['Intercept'])) and (beta0 <= (beta_hat['Intercept'] +
cv_005*bse['Intercept']))
    beta1_result_005[i] = (beta1 >= (beta_hat['x1'] - cv_005*bse['x1'])) and
(beta1 <= (beta_hat['x1'] + cv_005*bse['x1']))
    beta2_result_005[i] = (beta2 >= (beta_hat['x2'] - cv_005*bse['x2'])) and
(beta2 <= (beta_hat['x2'] + cv_005*bse['x2']))

    ## Whether the true parameters lie in the 99% CI's?
    beta0_result_001[i] = (beta0 >= (beta_hat['Intercept'] -
cv_001*bse['Intercept'])) and (beta0 <= (beta_hat['Intercept'] +
cv_001*bse['Intercept']))
    beta1_result_001[i] = (beta1 >= (beta_hat['x1'] - cv_001*bse['x1'])) and
(beta1 <= (beta_hat['x1'] + cv_001*bse['x1']))

```

```

    beta2_result_001[i] = (beta2 >= (beta_hat['x2'] - cv_001*bse['x2'])) and
    ↪(beta2 <= (beta_hat['x2'] + cv_001*bse['x2']))

```

```

## Print out the results
## Alpha = 0.05
print(np.mean(beta0_result_005), np.mean(beta1_result_005), np.
    ↪mean(beta2_result_005))
## Alpha = 0.01
print(np.mean(beta0_result_001), np.mean(beta1_result_001), np.
    ↪mean(beta2_result_001))

```

0.883 0.879 0.899

0.98 0.979 0.978

1.0.3 Q2

a.

```

[ ]: import wooldridge as woo
import statsmodels.formula.api as smf
import scipy.stats as stats
import numpy as np

## Import data
vote1 = woo.data('vote1')

## a. Descriptive statistics
vote1[['voteA', 'lexpendA', 'lexpendB', 'prtystrA']].describe()

```

```

[ ]:
      voteA    lexpendA    lexpendB    prtystrA
count  173.000000  173.000000  173.000000  173.000000
mean    50.502890    5.025556    4.944369   49.757225
std     16.784761    1.601602    1.571143    9.983650
min     16.000000   -1.197328   -0.072571   22.000000
25%     36.000000    4.402246    4.095244   44.000000
50%     50.000000    5.492164    5.400558   50.000000
75%     65.000000    6.125580    6.110837   56.000000
max     84.000000    7.293476    7.344844   71.000000

```

b.

β_1 is the percentage point increase in *voteA* for a one unit increase in *lexpendA*, holding other variables constant.

c.

$$H_0 : \beta_1 + \beta_2 = 0$$

d.

```
[ ]: ## d. Fit the regression
reg = smf.ols(formula = 'voteA ~ lexpendA + lexpendB + prtystA', data = vote1)
results = reg.fit()

print(f'results.summary():\n{results.summary()}\n')
```

results.summary():

```

                                OLS Regression Results
=====
Dep. Variable:                  voteA      R-squared:                0.793
Model:                            OLS      Adj. R-squared:            0.789
Method:                 Least Squares      F-statistic:                215.2
Date:                Tue, 04 Jun 2024      Prob (F-statistic):        1.76e-57
Time:                09:55:07      Log-Likelihood:            -596.86
No. Observations:                173      AIC:                        1202.
Df Residuals:                    169      BIC:                        1214.
Df Model:                          3
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	45.0789	3.926	11.481	0.000	37.328	52.830
lexpendA	6.0833	0.382	15.919	0.000	5.329	6.838
lexpendB	-6.6154	0.379	-17.463	0.000	-7.363	-5.868
prtystA	0.1520	0.062	2.450	0.015	0.030	0.274

```

=====
Omnibus:                        8.900      Durbin-Watson:              1.604
Prob(Omnibus):                  0.012      Jarque-Bera (JB):           8.832
Skew:                          0.493      Prob(JB):                   0.0121
Kurtosis:                      3.505      Cond. No.                   344.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

From the regression output, we cannot reject the null hypothesis that $\beta_1 + \beta_2 = 0$ at 5% significance level.

```
[ ]: ## Test H0: beta1 + beta 2 = 0 in b.
## Use method f_test
hypotheses = ['lexpendA + lexpendB = 0']
fctest = results.f_test(hypotheses)
fstat = fctest.fvalue          ## Extract the F statistic
fpval = fctest.pvalue          ## Extract the p-value

print(f'fstat: {round(fstat, 3)}\n')
```

```
print(f'fpval: {round(fpval, 3)}\n')
```

fstat: 0.996

fpval: 0.32

The p-value is larger than 0.05, so we cannot reject the null hypothesis at 5% significance level.
e.

```
[ ]: ## e. Estimate a modified model:
vote1['lB_lA'] = vote1['lexpendB'] - vote1['lexpendA']
reg_mod = smf.ols(formula = 'voteA ~ lB_lA + prtystA', data = vote1)
results_mod = reg_mod.fit()
print(f'results_mod.summary():\n{results_mod.summary()}\n')

## From the result, p-value of the t statistic (coefficient of lexpendA) is
↳ also about 0.32 and we cannot reject the null
## at 5% significant level. We do not have enough evidence to say that there is
↳ no offset effect.

## Verify whether  $t^2 = F$ 
tvalues = results_mod.tvalues
tvalues_sq = tvalues['lB_lA']**2
print(round(tvalues_sq,4), round(fstat,4))
```

results_mod.summary():

```

                                OLS Regression Results
=====
Dep. Variable:                  voteA      R-squared:                0.791
Model:                            OLS      Adj. R-squared:            0.789
Method:                 Least Squares      F-statistic:                322.3
Date:                Tue, 04 Jun 2024      Prob (F-statistic):          1.42e-58
Time:                09:55:07      Log-Likelihood:              -597.37
No. Observations:                173      AIC:                        1201.
Df Residuals:                    170      BIC:                        1210.
Df Model:                          2
Covariance Type:                  nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	42.7028	3.122	13.677	0.000	36.539	48.866
lB_lA	-6.3517	0.272	-23.394	0.000	-6.888	-5.816
prtystA	0.1464	0.062	2.370	0.019	0.024	0.268

```

=====
Omnibus:                        11.793      Durbin-Watson:              1.599
Prob(Omnibus):                  0.003      Jarque-Bera (JB):           12.607
Skew:                          0.554      Prob(JB):                   0.00183

```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

547.2726 0.9963

f.

```
[ ]: ## f. Test H0: beta1 + beta2 = 1.5
## Use the method f_test
hypotheses = ['lexpendA + lexpendB = 1.5']
fctest = results.f_test(hypotheses)
fstat = fctest.fvalue
fpval = fctest.pvalue

print(f'fstat: {round(fstat, 3)}\n')
print(f'fpval: {round(fpval, 3)}\n')

## Estimate a modified model
vote1['lA_lB'] = vote1['lexpendA'] - vote1['lexpendB']
reg_mod = smf.ols(formula = 'voteA ~ lA_lB + prtystA', data = vote1)
results_mod = reg_mod.fit()
print(f'results_mod.summary():\n{results_mod.summary()}\n')

## Manually calculate t statistic
n = vote1.shape[0]
k = 3
df = n - k - 1
b_lexpendA = results_mod.params['lA_lB']
se_lexpendA = results_mod.bse['lA_lB']
tstat = round((b_lexpendA - 1.5) / se_lexpendA, 3)
tpval = 2*stats.t.cdf(-abs(tstat), df = df)

print(f'tstat: {round(tstat, 3)}\n')
print(f'tpval: {round(tpval, 3)}\n') ## should be the same
    ↳ value as fpval

## Use method t_test
hypothesis = 'lA_lB = 1.5'
ttest = results_mod.t_test(hypothesis)
tstat = ttest.statistic[0][0]
tpval = ttest.pvalue

print(f'tstat: {round(tstat, 3)}\n')
```

```
print(f'tpval: {np.around(tpval, 3)}\n')          ## should be the same
↪value as fpval
```

fstat: 14.531

fpval: 0.0

results_mod.summary():

```

                                OLS Regression Results
=====
Dep. Variable:                  voteA      R-squared:                0.791
Model:                          OLS      Adj. R-squared:            0.789
Method:                        Least Squares  F-statistic:                322.3
Date:                          Tue, 04 Jun 2024  Prob (F-statistic):      1.42e-58
Time:                          09:55:07    Log-Likelihood:             -597.37
No. Observations:                173      AIC:                        1201.
Df Residuals:                    170      BIC:                        1210.
Df Model:                          2
Covariance Type:                  nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	42.7028	3.122	13.677	0.000	36.539	48.866
lA_lB	6.3517	0.272	23.394	0.000	5.816	6.888
prtystrA	0.1464	0.062	2.370	0.019	0.024	0.268

```

=====
Omnibus:                        11.793    Durbin-Watson:                1.599
Prob(Omnibus):                  0.003    Jarque-Bera (JB):              12.607
Skew:                          0.554    Prob(JB):                      0.00183
Kurtosis:                      3.721    Cond. No.                      270.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

tstat: 17.869

tpval: 0.0

tstat: 17.869

tpval: 0.0

1.0.4 Q3

a.

```
[ ]: import wooldridge as woo
import numpy as np
import statsmodels.formula.api as smf
import scipy.stats as stats

## Import data
hprice1 = woo.data('hprice1')

## a. Descriptive statistics
hprice1[['price', 'sqrft', 'bdrms']].describe()
```

```
[ ]:      price      sqrft      bdrms
count  88.000000   88.000000  88.000000
mean   293.546034  2013.693182   3.568182
std    102.713445   577.191583   0.841393
min    111.000000  1171.000000   2.000000
25%    230.000000  1660.500000   3.000000
50%    265.500000  1845.000000   3.000000
75%    326.250000  2227.000000   4.000000
max    725.000000  3880.000000   7.000000
```

b.

```
[ ]: ## b. Fit the regression
reg = smf.ols(formula = 'np.log(price) ~ sqrft + bdrms', data = hprice1)
results = reg.fit()

## Estimate theta1: 150 sqrft bed room added
theta1 = 150*results.params['sqrft'] + results.params['bdrms']
theta1 = round(theta1, 4)
print(f'Percentage change:\n{theta1*100}%\n')
```

Percentage change:
8.58%

c.

```
[ ]: ## c.
## Calculate standard error of theta1 and construct CI's
hprice1['new_var'] = hprice1['sqrft'] - 150*hprice1['bdrms']
reg_mod = smf.ols(formula = 'np.log(price) ~ new_var + bdrms', data = hprice1)
results_mod = reg_mod.fit()
print(round(results_mod.bse['bdrms'], 4))
```

0.0268

standard error of the percentage change in b is 0.0268

```
[ ]: ## Now the slope parameters of bdrms is theta1
bse = results_mod.bse['bdrms']
n = hprice1.shape[0]
k = 2
df = n - k - 1
cv_005 = stats.t.ppf(1 - 0.025, df = df) ## Critical value for 95% CI
cv_001 = stats.t.ppf(1 - 0.005, df = df) ## Critical value for 99% CI

ci_005_up = theta1 + cv_005*bse
ci_005_low = theta1 - cv_005*bse
ci_005 = [round(ci_005_low, 3), round(ci_005_up, 3)]
print(f'ci_005:\n{ci_005}\n') ## 95% CI for theta1

ci_001_up = theta1 + cv_001*bse
ci_001_low = theta1 - cv_001*bse
ci_001 = [round(ci_001_low, 3), round(ci_001_up, 3)]
print(f'ci_001:\n{ci_001}\n') ## 99% CI for theta1

ci_005:
[0.033, 0.139]

ci_001:
[0.015, 0.156]
```

Based on the confidence intervals do not contain zero. This provides statistical evidence that adding a 150-square-foot bedroom significantly increases the house price, with the percentage increase ranging from approximately 3.3% to 15.6% (95% confidence interval).